

Chapter 2: Intelligent agents and problem solving


Artificial Intelligence and Machine Learning

Intelligent Agents

- ◇ Agents and environments
- ◇ Rationality
- ◇ PEAS (Performance measure, Environment, Actuators, Sensors)
- ◇ Environment types
- ◇ Agent types

Agents and Environments

- An **Agent** is anything that can be viewed as perceiving its **environment** through **sensors** and acting upon that environment through **actuators**




Agents and Environments

- Percept – the agent's perceptual inputs
 - percept sequence is a sequence of everything the agent has ever perceived
- Agent Function – describes the agent's behavior
 - Maps any given percept sequence to an action
 - $f : P^* \rightarrow A$
- Agent Program – an implementation of an agent function for an artificial agent

Agents and Environments

- Example: Vacuum Cleaner World
 - Two locations: squares A and B
 - Perceives what square it is in
 - Perceives if there is dirt in the current square
 - Actions
 - move left
 - move right
 - suck up the dirt
 - do nothing



Agents and Environments

- Agent Function:
Vacuum Cleaner World
 - If the current square is dirty, then suck, otherwise move to the other square

Percept Sequence	Action
[A, Clean]	Right
[A, Dirty]	Suck
[B, Clean]	Left
[B, Dirty]	Suck
[A, Clean], [A, Clean]	Right
[A, Clean], [A, Dirty]	Suck

Agents and Environments

- But what is the right way to fill out the table?
 - is the agent
 - good or bad
 - intelligent or stupid
 - can it be implemented in a small program?

```
Function Reflex-Vacuum-Agent([location, status]) return an action
  if status == Dirty then return Suck
  else if location = A then return Right
  else if location = B then return Left
```

Good Behavior and Rationality

- Rational Agent – an agent that does the “right” thing
 - Every entry in the table for the agent function is filled out correctly
 - Doing the right thing is better than doing the wrong thing
 - What does it mean to do the right thing?

Good Behavior and Rationality

- Performance Measure
 - A scoring function for evaluating the environment space
- **Rational Agent** – for each possible percept sequence, a rational agent should select an action that is expected to maximize its performance measure, given the evidence provided by the percept sequence and what ever built-in knowledge the agent has.

Good Behavior and Rationality

- Rational \neq omniscient
- Rational \neq clairvoyant
- Rational \neq successful

- Rational \rightarrow exploration, learning, autonomy

The Nature of Environments

- Task environments
 - The “problems” to which a rational agent is the “solution”
- PEAS
 - Performance
 - Environment
 - Actuators
 - Sensors

The Nature of Environments

- Properties of task environments
 - Fully Observable vs. Partially Observable
 - Deterministic vs. Stochastic
 - Episodic vs. Sequential
 - Static vs. Dynamic
 - Discrete vs. Continuous
 - Single agent vs. Multi-agent
- The real world is partially observable, stochastic, sequential, dynamic, continuous, multi-agent

The Nature of Environments

- Agent Examples
 - Automated Taxi
 - Mars Rover
 - Trader
 - Chess

The Structure of Agents

- Agent = Architecture + Program
- Basic algorithm for a rational agent
 - While (true) do
 - Get percept from sensors into memory
 - Determine best action based on memory
 - Record action in memory
 - Perform action
- Most AI programs are a variation of this theme

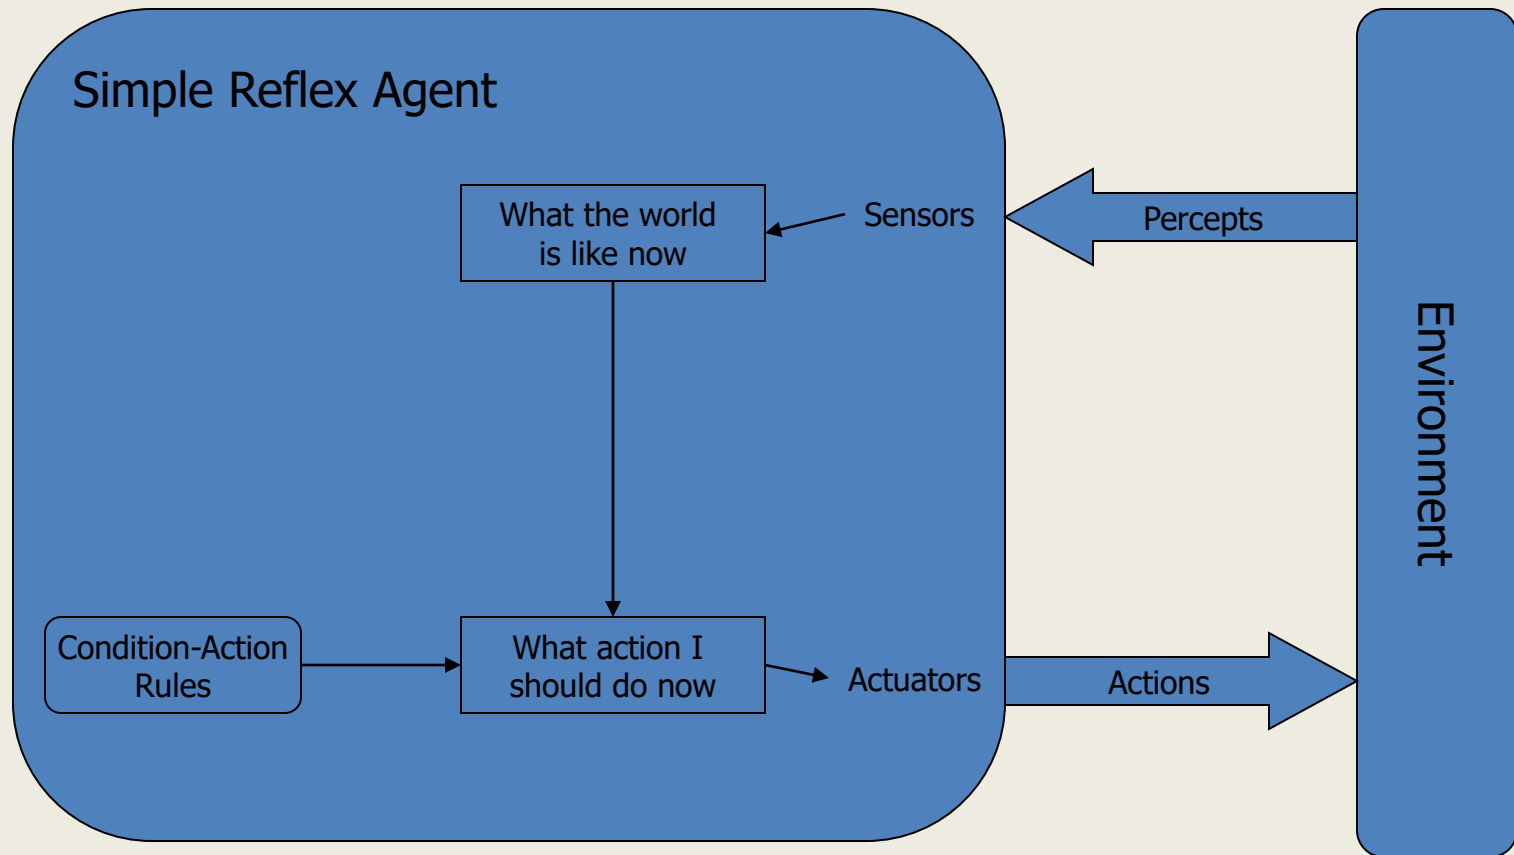
The Structure of Agents

- Table Driven Agent

```
function Table-Driven-Agent (percept) return action
static:      percepts, a sequence, initially empty
              table, a table of actions, indexed by
              percept sequences, initially fully
              specified

append percept to the end of the table
action <- LOOKUP( percept, table )
return action
```

The Structure of Agents



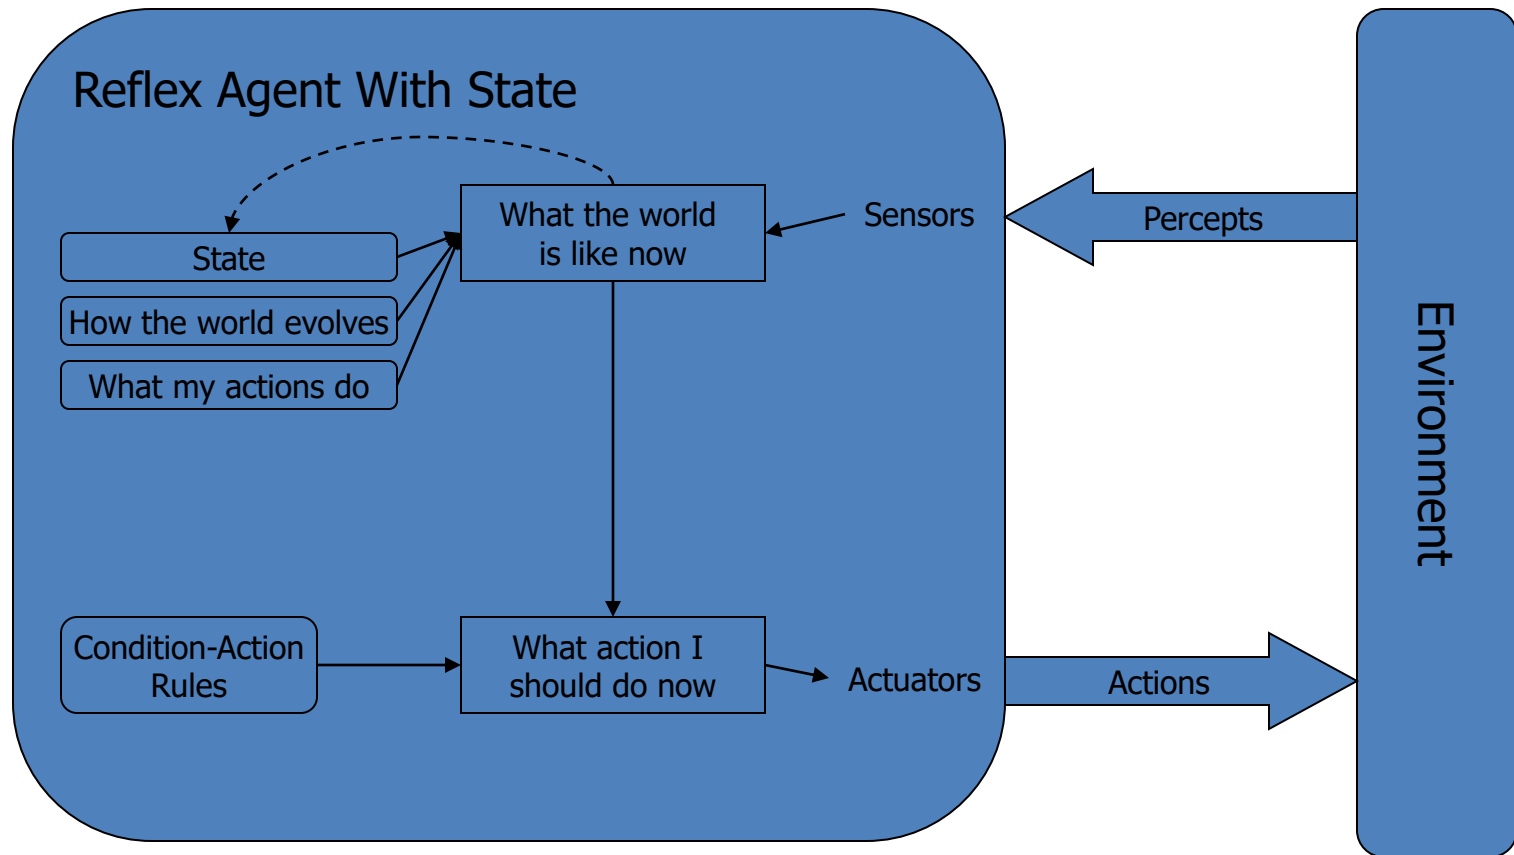
The Structure of Agents

- Simple Reflex Agent

```
function Simple-Reflex-Agent (percept) return action
static:      rules, a set of condition-action rules

    state <- INTERPRET-INPUT( percept )
    rule <- RULE-MATCH( state, rules )
    action <- RULE-ACTION[ rule ]
    return action
```

The Structure of Agents



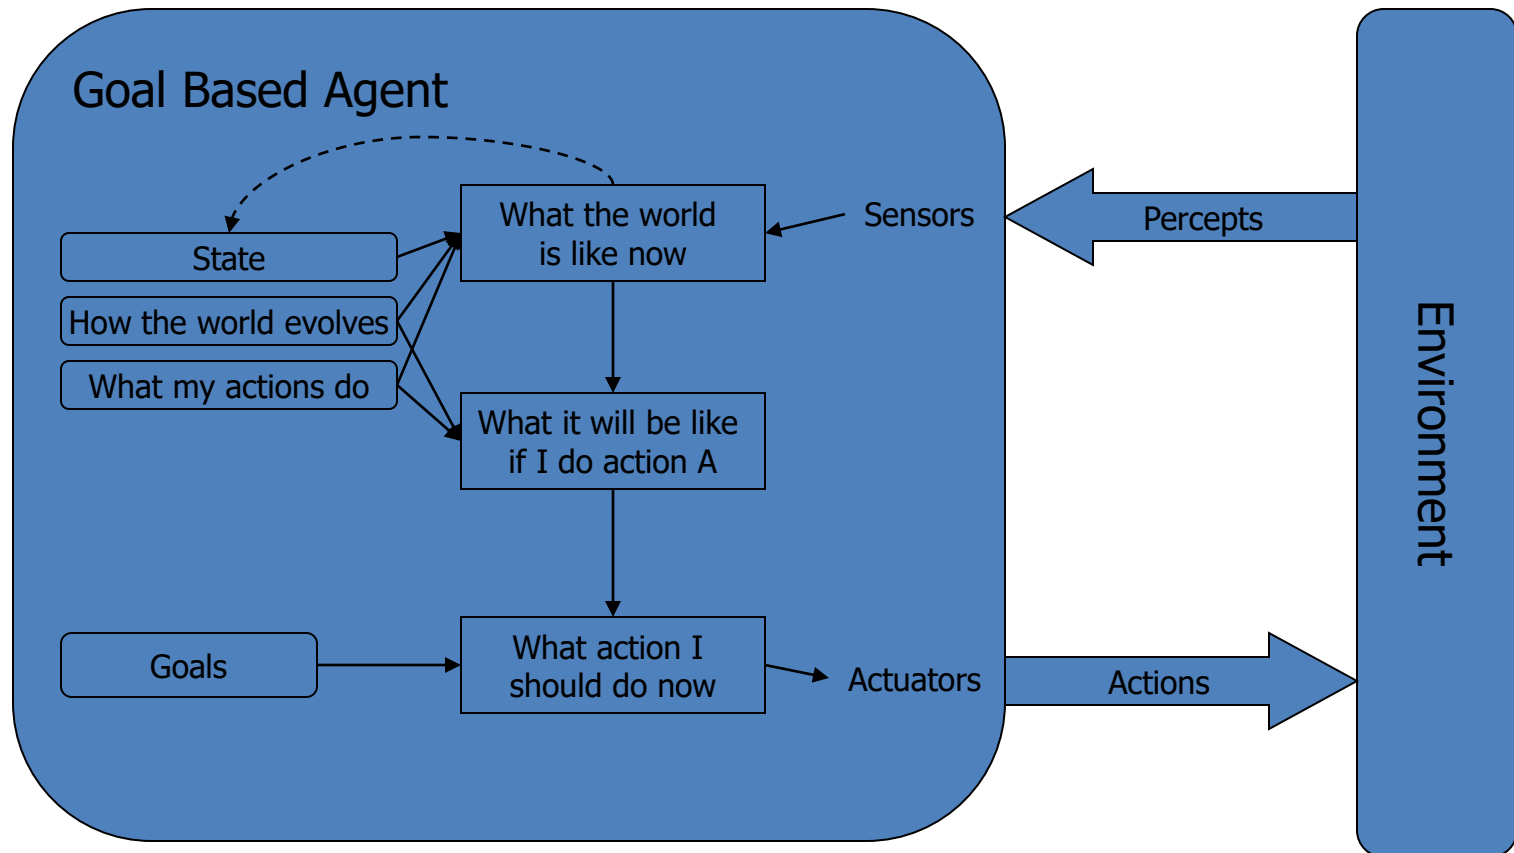
The Structure of Agents

- Reflex Agent With State

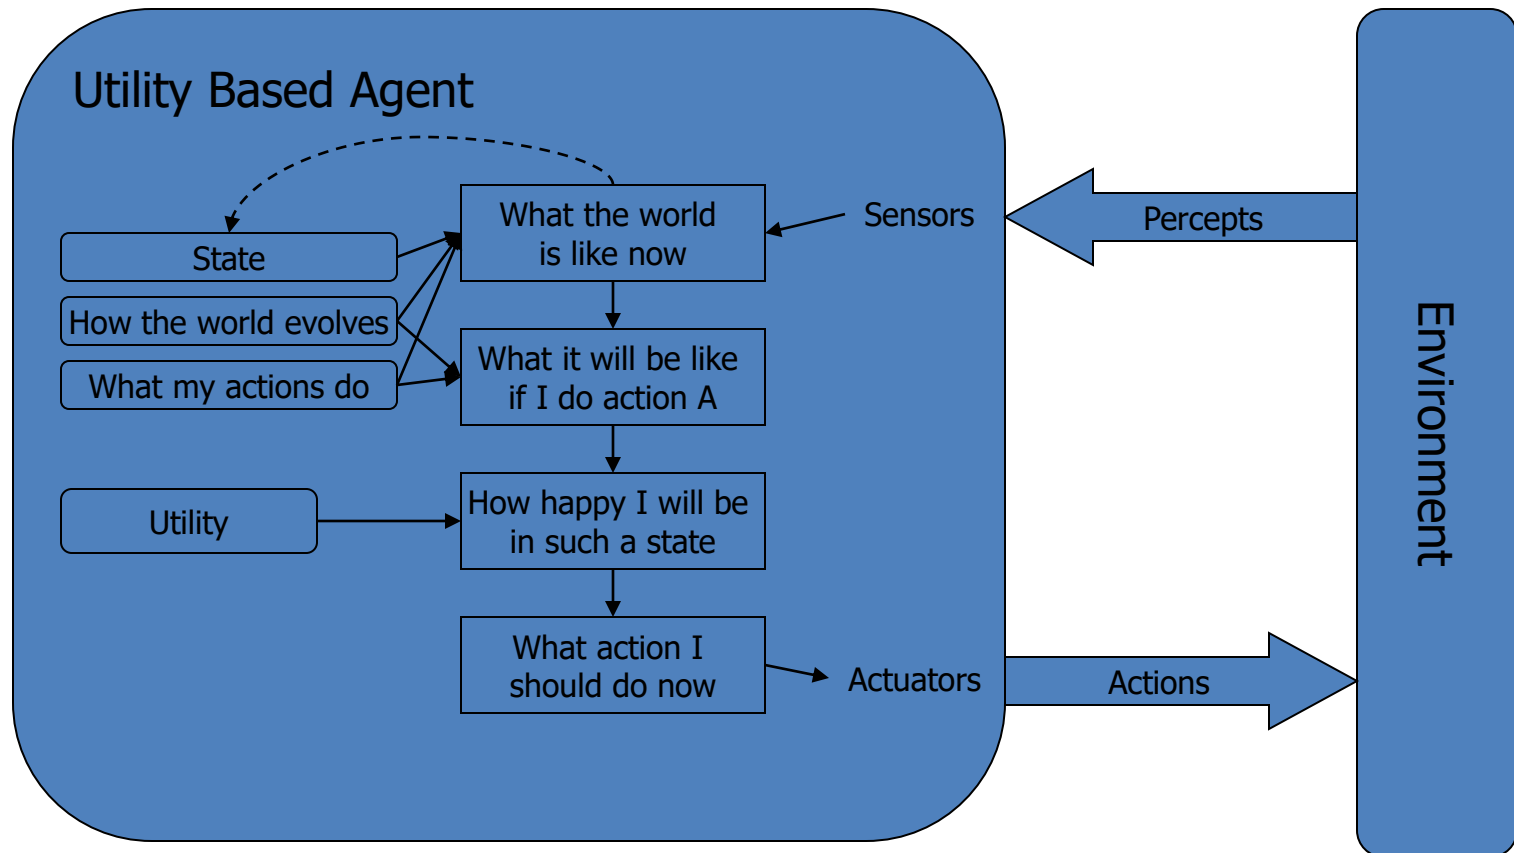
```
function Reflex-Agent-With-State (percept) return action
static: state, a description of the current world state
        rules, a set of condition-action rules
        action, the most recent action, initially none
```

```
state <- UPDATE-STATE( state, action, percept )
rule <- RULE-MATCH( state, rules )
action <- RULE-ACTION[ rule ]
return action
```

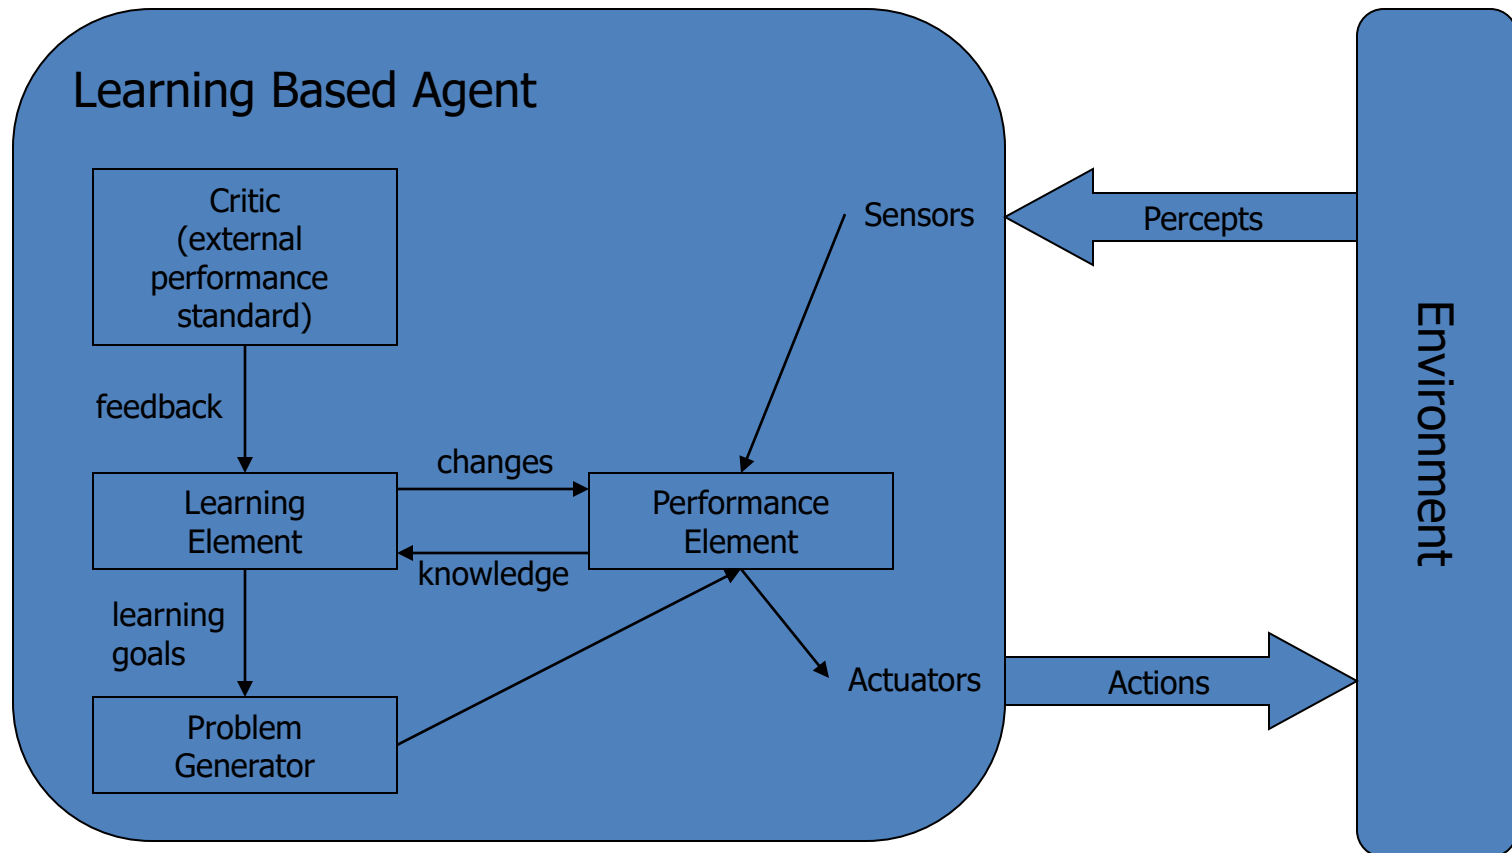
The Structure of Agents



The Structure of Agents



The Structure of Agents



Problem Solving and Search

Problem solving

Problem types

Problem formulation


Example problems

Basic search algorithms

Problem Solving Agents

- Problem solving agent
 - A kind of “goal based” agent
 - Finds sequences of actions that lead to desirable states.
- The algorithms are uninformed
 - No extra information about the problem other than the definition
 - No extra information
 - No heuristics (rules)

Goal Based Agent



Goal Based Agent

Function Simple-Problem-Solving-Agent(percept) returns action

Inputs: percept	a percept
Static: seq	an action sequence initially empty
state	some description of the current world
goal	a goal, initially null
problem	a problem formulation

```
state <- UPDATE-STATE( state, percept )
if seq is empty then do
  goal <- FORMULATE-GOAL( state )
  problem <- FORMULATE-PROBLEM( state, goal )
  seq <- SEARCH( problem )           # SEARCH
action <- RECOMMENDATION ( seq )     # SOLUTION
seq <- REMAINDER( seq )
return action                       # EXECUTION
```

Goal Based Agents

- Assumes the problem environment is:
 - Static
 - The plan remains the same
 - Observable
 - Agent knows the initial state
 - Discrete
 - Agent can enumerate the choices
 - Deterministic
 - Agent can plan a sequence of actions such that each will lead to an intermediate state
- The agent carries out its plans with its eyes closed
 - Certain of what's going on
 - Open loop system

Well Defined Problems and Solutions

- A problem
 - Initial state
 - Actions and Successor Function
 - Goal test
 - Path cost

Example: Water Pouring

- Given a 4 gallon bucket and a 3 gallon bucket, how can we measure exactly 2 gallons into one bucket?
 - There are no markings on the bucket
 - You must fill each bucket completely

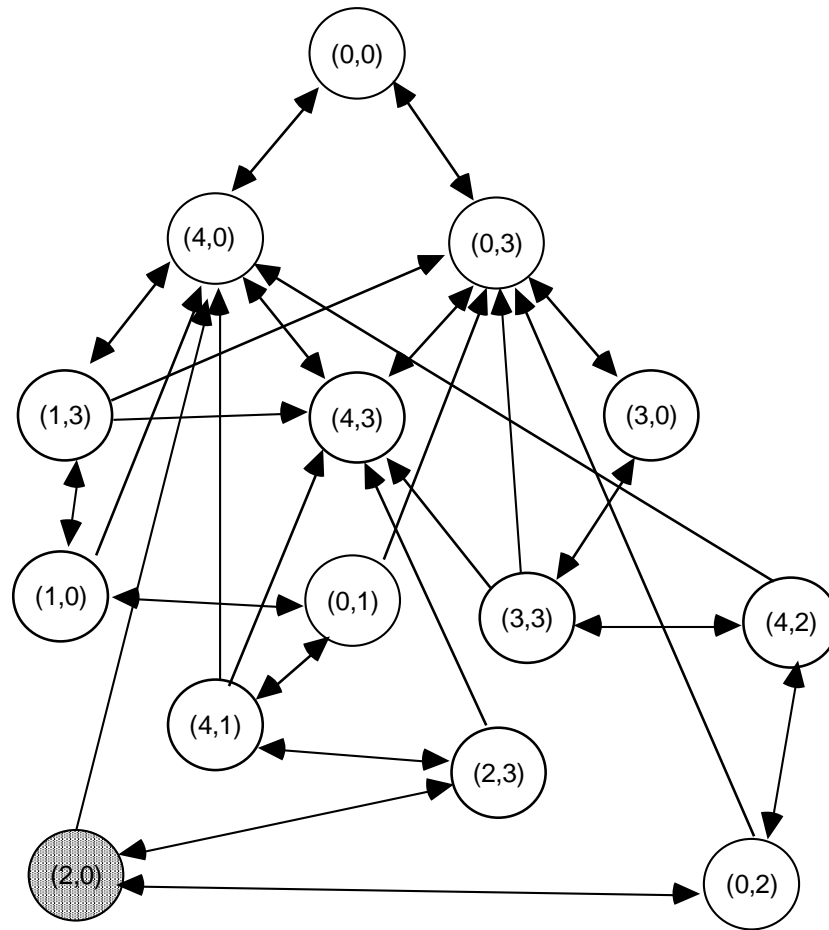
Example: Water Pouring

- Initial state:
 - The buckets are empty
 - Represented by the tuple (0 0)
- Goal state:
 - One of the buckets has two gallons of water in it
 - Represented by either (x 2) or (2 x)
- Path cost:
 - 1 per unit step

Example: Water Pouring

- Actions and Successor Function
 - Fill a bucket
 - $(x\ y) \rightarrow (3\ y)$
 - $(x\ y) \rightarrow (x\ 4)$
 - Empty a bucket
 - $(x\ y) \rightarrow (0\ y)$
 - $(x\ y) \rightarrow (x\ 0)$
 - Pour contents of one bucket into another
 - $(x\ y) \rightarrow (0\ x+y)$ or $(x+y-4,\ 4)$
 - $(x\ y) \rightarrow (x+y\ 0)$ or $(3,\ x+y-3)$

Example: Water Pouring



Example: Eight Puzzle

- States:
 - Description of the eight tiles and location of the blank tile
- Successor Function:
 - Generates the legal states from trying the four actions {*Left, Right, Up, Down*}
- Goal Test:
 - Checks whether the state matches the goal configuration
- Path Cost:
 - Each step costs 1

7	2	4
5		6
8	3	1


1	2	3
4	5	6
7	8	

Example: Eight Puzzle

- Eight puzzle is from a family of “sliding –block puzzles”
 - NP Complete
 - 8 puzzle has $9!/2 = 181440$ states
 - 15 puzzle has approx. $1.3 \cdot 10^{12}$ states
 - 24 puzzle has approx. $1 \cdot 10^{25}$ states


Example: Eight Queens

- Place eight queens on a chess board such that no queen can attack another queen
- No path cost because only the final state counts!
- Incremental formulations
- Complete state formulations




Example: Eight Queens

- States:
 - Any arrangement of 0 to 8 queens on the board
- Initial state:
 - No queens on the board
- Successor function:
 - Add a queen to an empty square
- Goal Test:
 - 8 queens on the board and none are attacked
- $64 * 63 * \dots * 57 = 1.8 * 10^{14}$ possible sequences
 - Ouch!

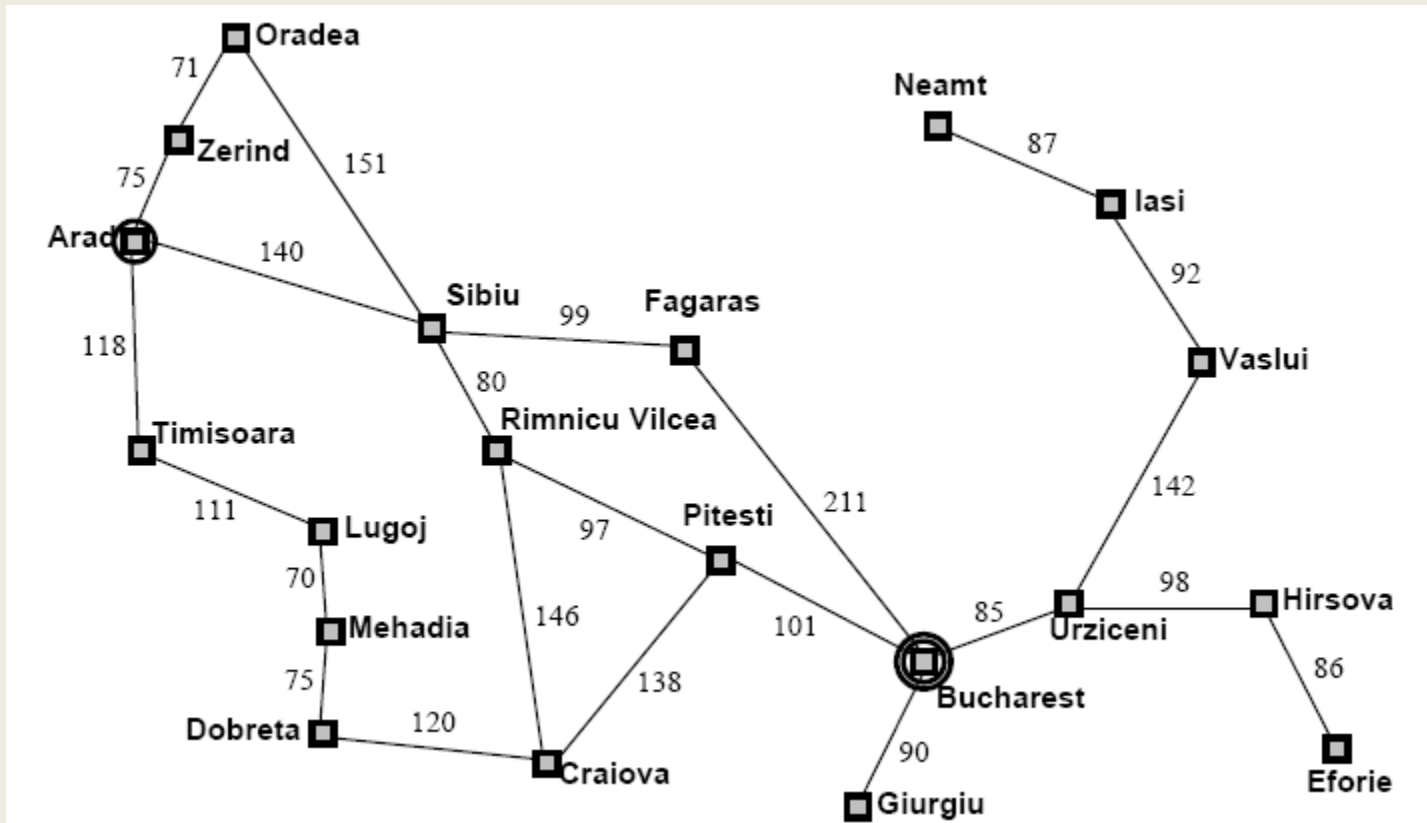


Example: Eight Queens

- States:
 - Arrangements of n queens, one per column in the leftmost n columns, with no queen attacking another are states
- Successor function:
 - Add a queen to any square in the leftmost empty column such that it is not attacked by any other queen.
- 2057 sequences to investigate



Example: Map Planning



Searching For Solutions

- Initial State
 - e.g. “At Arad”
- Successor Function
 - A set of action state pairs
 - $S(\text{Arad}) = \{(\text{Arad} \rightarrow \text{Zerind}, \text{Zerind}), \dots\}$
- Goal Test
 - e.g. $x = \text{“at Bucharest”}$
- Path Cost
 - sum of the distances traveled

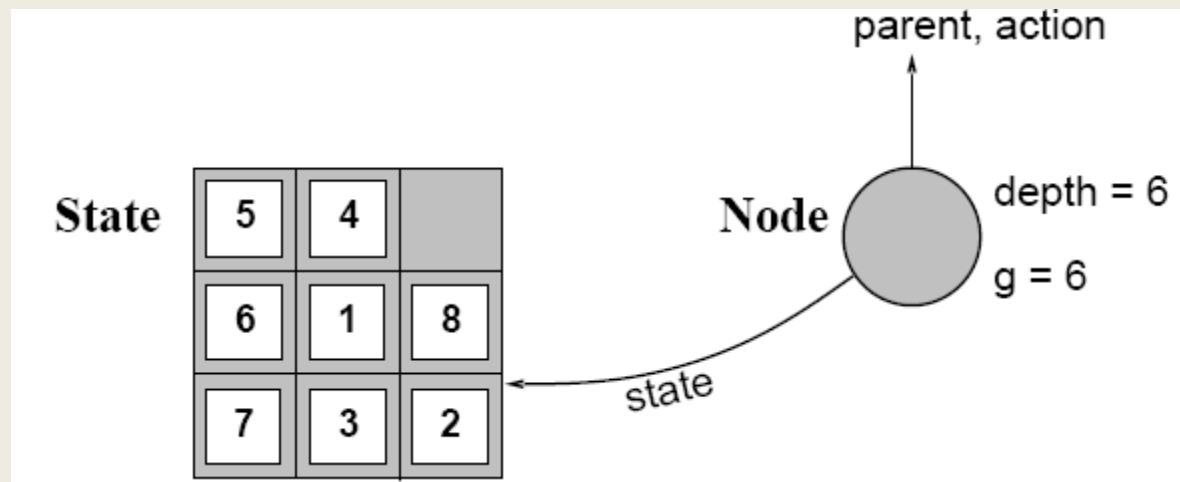
Searching For Solutions

- Having formulated some problems...how do we solve them?
- Search through a state space
- Use a search tree that is generated with an initial state and successor functions that define the state space

Searching For Solutions

- A **state** is (a representation of) a physical configuration
- A **node** is a data structure constituting part of a search tree
 - Includes parent, children, depth, path cost
- States do not have children, depth, or path cost
- The EXPAND function creates new nodes, filling in the various fields and using the SUCCESSOR function of the problem to create the corresponding states

Searching For Solutions



Searching For Solutions

```
function TREE-SEARCH( problem, strategy) returns a solution, or failure
  initialize the search tree using the initial state of problem
  loop do
    if there are no candidates for expansion then return failure
    choose a leaf node for expansion according to strategy
    if the node contains a goal state then return the corresponding solution
    else expand the node and add the resulting nodes to the search tree
  end
```

Searching For Solutions

```
function TREE-SEARCH( problem, fringe) returns a solution, or failure
  fringe  $\leftarrow$  INSERT(MAKE-NODE(INITIAL-STATE[problem]), fringe)
  loop do
    if fringe is empty then return failure
    node  $\leftarrow$  REMOVE-FRONT(fringe)
    if GOAL-TEST[problem] applied to STATE(node) succeeds return node
    fringe  $\leftarrow$  INSERTALL(EXPAND(node, problem), fringe)
```

```
function EXPAND( node, problem) returns a set of nodes
  successors  $\leftarrow$  the empty set
  for each action, result in SUCCESSOR-FN[problem](STATE[node]) do
    s  $\leftarrow$  a new NODE
    PARENT-NODE[s]  $\leftarrow$  node; ACTION[s]  $\leftarrow$  action; STATE[s]  $\leftarrow$  result
    PATH-COST[s]  $\leftarrow$  PATH-COST[node] + STEP-COST(node, action, s)
    DEPTH[s]  $\leftarrow$  DEPTH[node] + 1
    add s to successors
  return successors
```

Uninformed Search Strategies

- **Uninformed** strategies use only the information available in the problem definition
 - Also known as blind searching
- Breadth-first search
- Uniform-cost search
- Depth-first search
- Depth-limited search
- Iterative deepening search

Comparing Uninformed Search Strategies

- Completeness
 - Will a solution always be found if one exists?
- Time
 - How long does it take to find the solution?
 - Often represented as the number of nodes searched
- Space
 - How much memory is needed to perform the search?
 - Often represented as the maximum number of nodes stored at once
- Optimal
 - Will the optimal (least cost) solution be found?
- Page 81 in AIMA text


Comparing Uninformed Search Strategies

- Time and space complexity are measured in
 - b – maximum branching factor of the search tree
 - m – maximum depth of the state space
 - d – depth of the least cost solution


Breadth-First Search

- Recall from Data Structures the basic algorithm for a breadth-first search on a graph or tree
- Expand the **shallowest** unexpanded node
- Place all new successors at the end of a FIFO queue


Breadth-First Search




Breadth-First Search



Breadth-First Search



Breadth-First Search



Properties of Breadth-First Search

- Complete
 - Yes if b (max branching factor) is finite
- Time
 - $1 + b + b^2 + \dots + b^d + b(b^d - 1) = O(b^{d+1})$
 - exponential in d
- Space
 - $O(b^{d+1})$
 - Keeps every node in memory
 - This is the big problem; an agent that generates nodes at 10 MB/sec will produce 860 MB in 24 hours
- Optimal
 - Yes (if cost is 1 per step); not optimal in general

Lessons From Breadth First Search

- The memory requirements are a bigger problem for breadth-first search than is execution time
- Exponential-complexity search problems cannot be solved by uniformed methods for any but the smallest instances

Uniform-Cost Search

- Same idea as the algorithm for breadth-first search...but...
 - Expand the **least-cost** unexpanded node
 - FIFO queue is ordered by cost
 - Equivalent to regular breadth-first search if all step costs are equal


Uniform-Cost Search

- Complete
 - Yes if the cost is greater than some threshold
 - step cost $\geq \epsilon$
- Time
 - Complexity cannot be determined easily by d or b
 - Let C^* be the cost of the optimal solution
 - $O(b^{\lceil C^* / \epsilon \rceil})$
- Space
 - $O(b^{\lceil C^* / \epsilon \rceil})$
- Optimal
 - Yes, Nodes are expanded in increasing order


Depth-First Search

- Recall from Data Structures the basic algorithm for a depth-first search on a graph or tree
- Expand the *deepest* unexpanded node
- Unexplored successors are placed on a stack until fully explored


Depth-First Search




Depth-First Search




Depth-First Search




Depth-First Search




Depth-First Search




Depth-First Search




Depth-First Search




Depth-First Search




Depth-First Search




Depth-First Search



Depth-First Search



Depth-First Search



Depth-First Search

- Complete
 - No: fails in infinite-depth spaces, spaces with loops
 - Modify to avoid repeated spaces along path
 - Yes: in finite spaces
- Time
 - $O(b^m)$
 - Not great if m is much larger than d
 - But if the solutions are dense, this may be faster than breadth-first search
- Space
 - $O(bm)$...linear space
- Optimal
 - No

Depth-Limited Search

- A variation of depth-first search that uses a depth limit
 - Alleviates the problem of unbounded trees
 - Search to a predetermined depth ℓ (“ell”)
 - Nodes at depth ℓ have no successors
- Same as depth-first search if $\ell = \infty$
- Can terminate for failure and cutoff

Depth-Limited Search

Recursive implementation:

```
function DEPTH-LIMITED-SEARCH(problem, limit) returns soln/fail/cutoff
  RECURSIVE-DLS(MAKE-NODE(INITIAL-STATE[problem]), problem, limit)

function RECURSIVE-DLS(node, problem, limit) returns soln/fail/cutoff
  cutoff-occurred?  $\leftarrow$  false
  if GOAL-TEST[problem](STATE[node]) then return node
  else if DEPTH[node] = limit then return cutoff
  else for each successor in EXPAND(node, problem) do
    result  $\leftarrow$  RECURSIVE-DLS(successor, problem, limit)
    if result = cutoff then cutoff-occurred?  $\leftarrow$  true
    else if result  $\neq$  failure then return result
  if cutoff-occurred? then return cutoff else return failure
```


Depth-Limited Search

- Complete
 - Yes if $\ell < d$
- Time
 - $O(b^\ell)$
- Space
 - $O(b\ell)$
- Optimal
 - No if $\ell > d$

Iterative Deepening Search

- Iterative deepening depth-first search
 - Uses depth-first search
 - Finds the best depth limit
 - Gradually increases the depth limit; 0, 1, 2, ... until a goal is found

Informed Search

Best-first search

A* search

Heuristics

Hill climbing

Iterative improvement algorithms

Informed (Heuristic) Search Strategies

- **Informed Search** – a strategy that uses problem-specific knowledge beyond the definition of the problem itself
- **Best-First Search** – an algorithm in which a node is selected for expansion based on an evaluation function $f(n)$
 - Traditionally the node with the lowest evaluation function is selected
 - Not an accurate name...expanding the best node first would be a straight march to the goal.
 - Choose the node that *appears* to be the best

Informed (Heuristic) Search Strategies

- There is a whole family of Best-First Search algorithms with different evaluation functions
 - Each has a heuristic function $h(n)$
- $h(n)$ = estimated cost of the cheapest path from node n to a goal node
- Example: in route planning the estimate of the cost of the cheapest path might be the straight line distance between two cities


A Quick Review

- $g(n)$ = cost from the initial state to the current state n
- $h(n)$ = estimated cost of the cheapest path from node n to a goal node
- $f(n)$ = evaluation function to select a node for expansion (usually the lowest cost node)

Greedy Best-First Search

- Greedy Best-First search tries to expand the node that is closest to the goal assuming it will lead to a solution quickly
 - $f(n) = h(n)$
 - aka “Greedy Search”
- Implementation
 - expand the “most desirable” node into the fringe queue
 - sort the queue in decreasing order of desirability
- Example: consider the straight-line distance heuristic h_{SLD}
 - Expand the node that appears to be closest to the goal

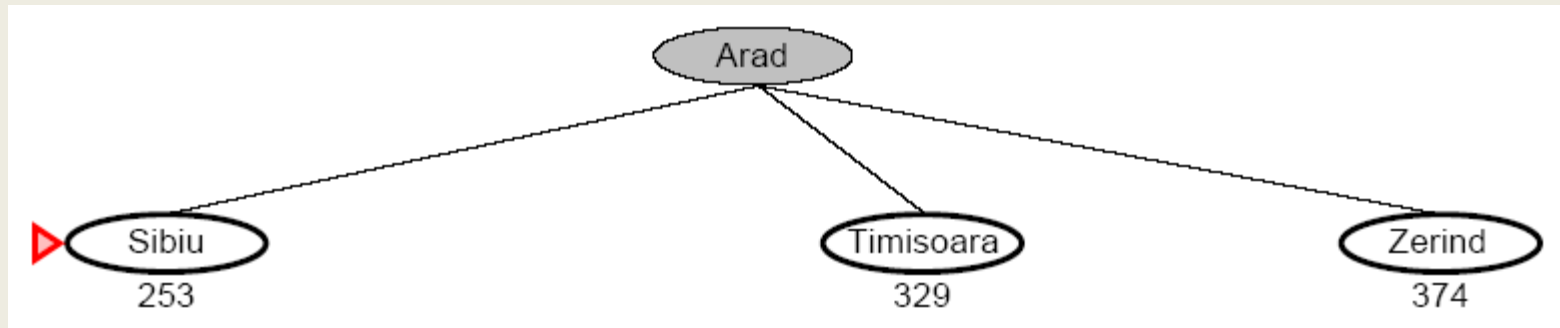
Greedy Best-First Search



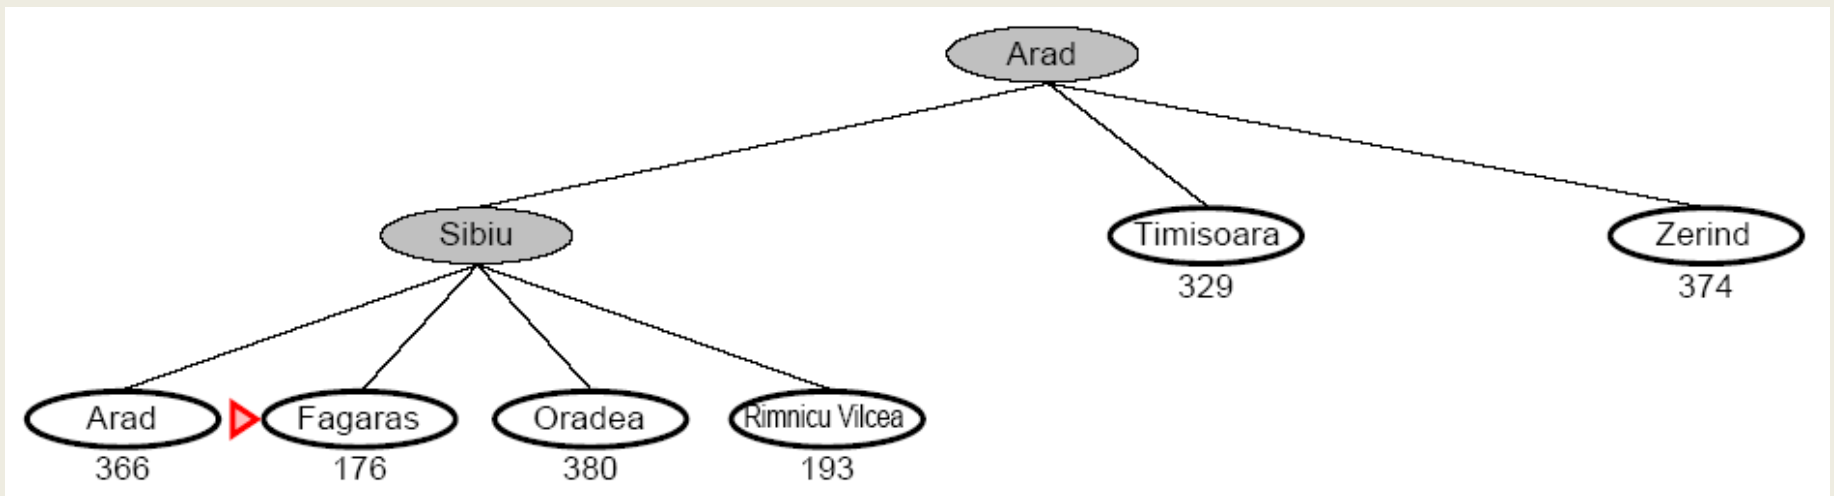
Greedy Best-First Search

- $h_{\text{SLD}}(\text{In}(\text{Arid})) = 366$
- Notice that the values of h_{SLD} cannot be computed from the problem itself
- It takes some experience to know that h_{SLD} is correlated with actual road distances
 - Therefore a useful heuristic

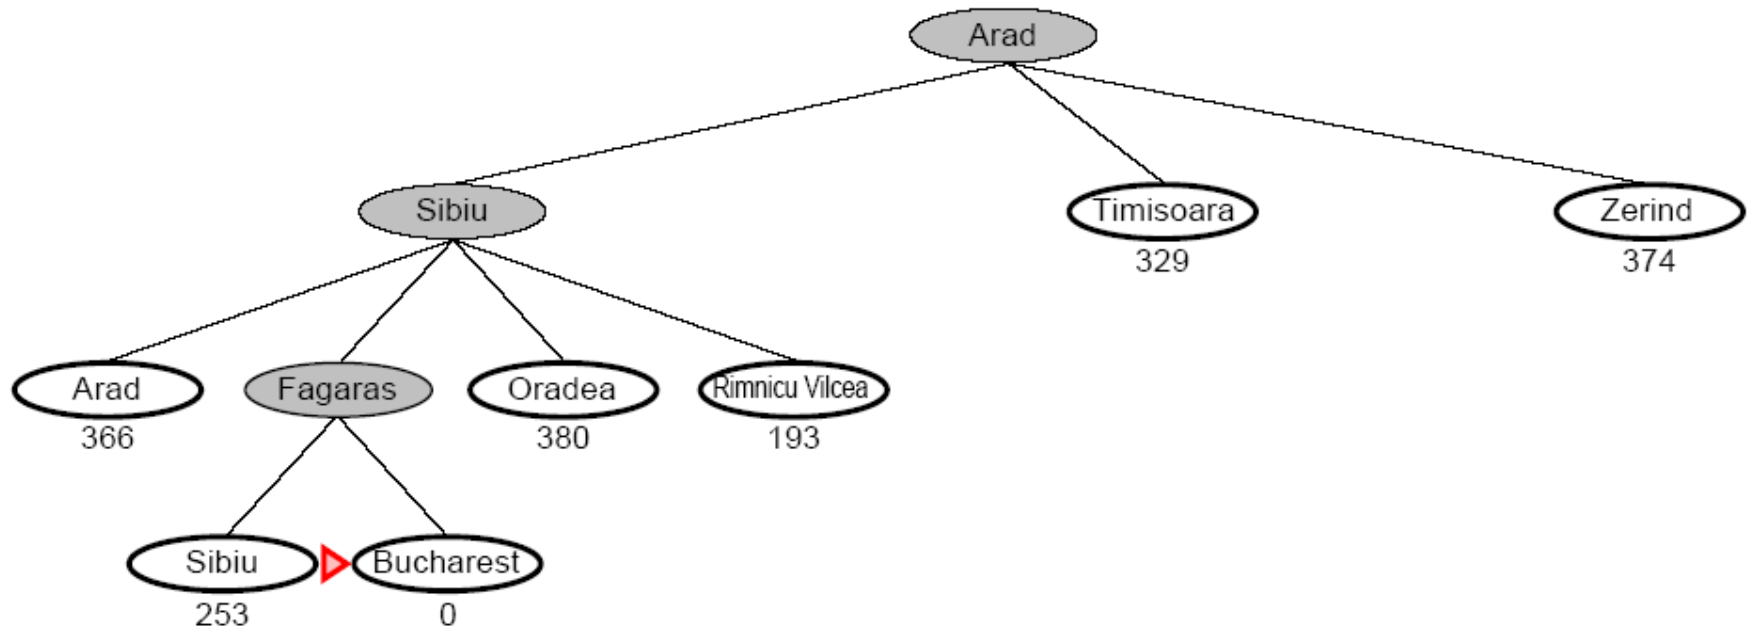
Greedy Best-First Search



Greedy Best-First Search



Greedy Best-First Search



Greedy Best-First Search

- Complete
 - No, GBFS can get stuck in loops (e.g. bouncing back and forth between cities)
- Time
 - $O(b^m)$ but a good heuristic can have dramatic improvement
- Space
 - $O(b^m)$ – keeps all the nodes in memory
- Optimal
 - No!

A Quick Review - Again

- $g(n)$ = cost from the initial state to the current state n
- $h(n)$ = estimated cost of the cheapest path from node n to a goal node
- $f(n)$ = evaluation function to select a node for expansion (usually the lowest cost node)

A* Search

- A* (A star) is the most widely known form of Best-First search
 - It evaluates nodes by combining $g(n)$ and $h(n)$
 - $f(n) = g(n) + h(n)$
 - Where
 - $g(n)$ = cost so far to reach n
 - $h(n)$ = estimated cost to goal from n
 - $f(n)$ = estimated total cost of path through n

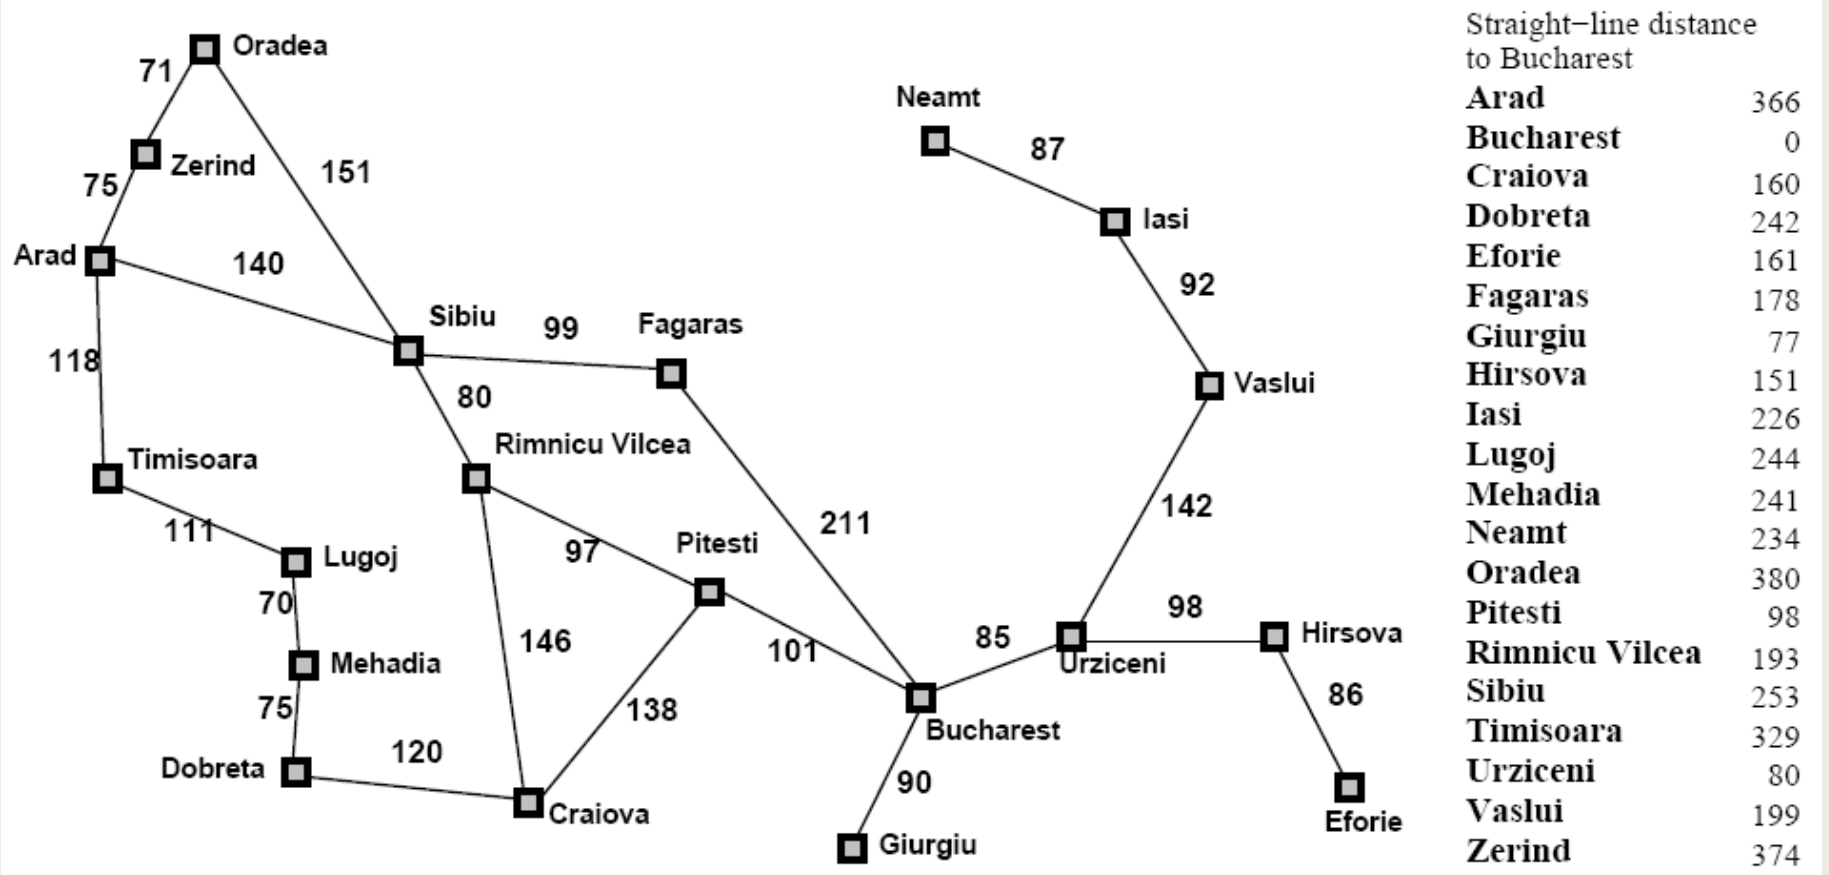
A* Search

- When $h(n) = \text{actual cost to goal}$
 - Only nodes in the correct path are expanded
 - Optimal solution is found
- When $h(n) < \text{actual cost to goal}$
 - Additional nodes are expanded
 - Optimal solution is found
- When $h(n) > \text{actual cost to goal}$
 - Optimal solution can be overlooked

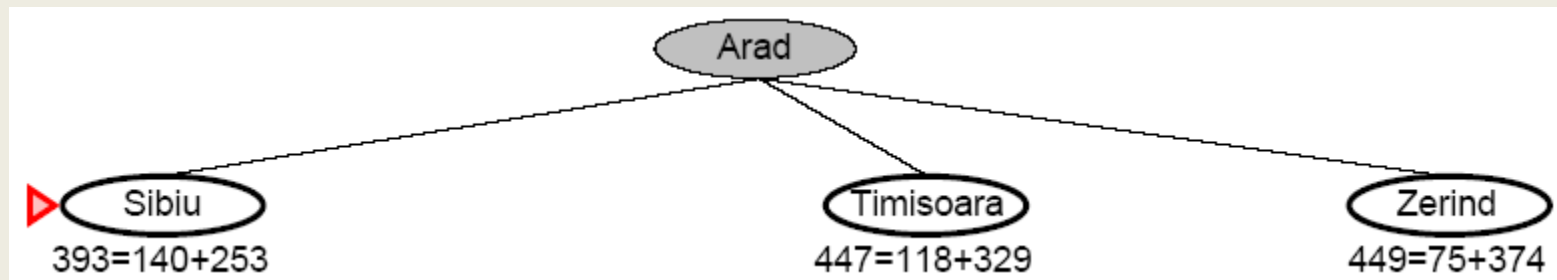

A* Search

- A* is optimal if it uses an ***admissible heuristic***
 - $h(n) \leq h^*(n)$ the true cost from node n
 - if $h(n)$ *never overestimates* the cost to reach the goal
- Example
 - h_{SLD} never overestimates the actual road distance


Greedy Best-First Search



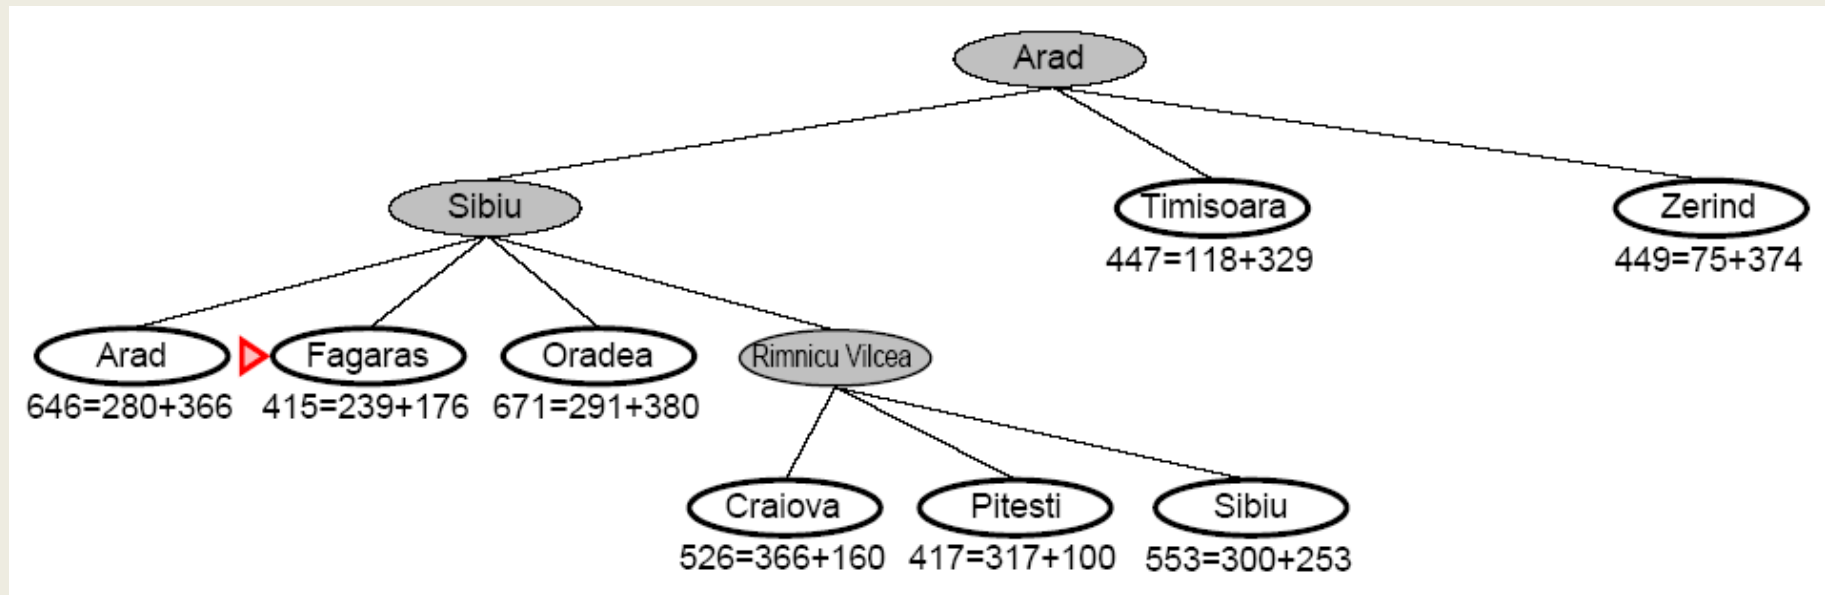
A* Search



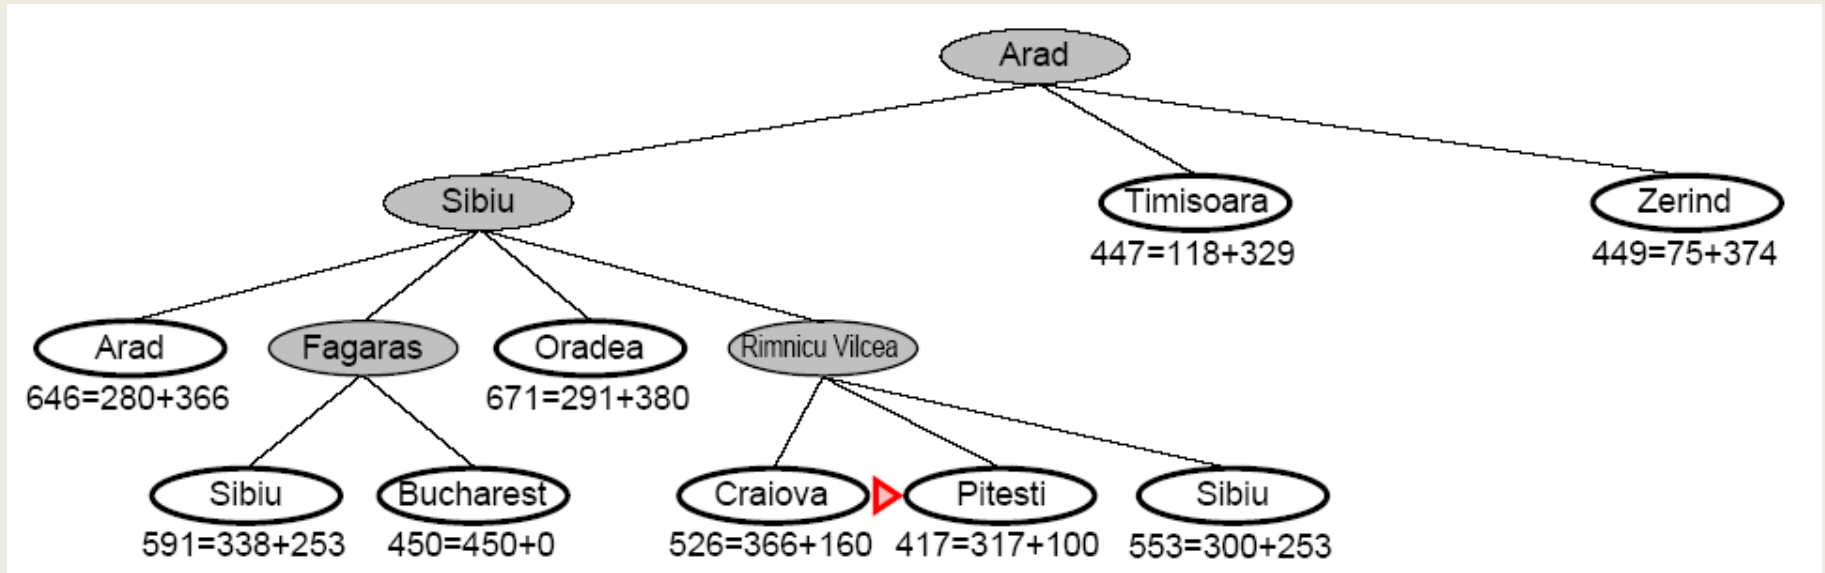
A* Search



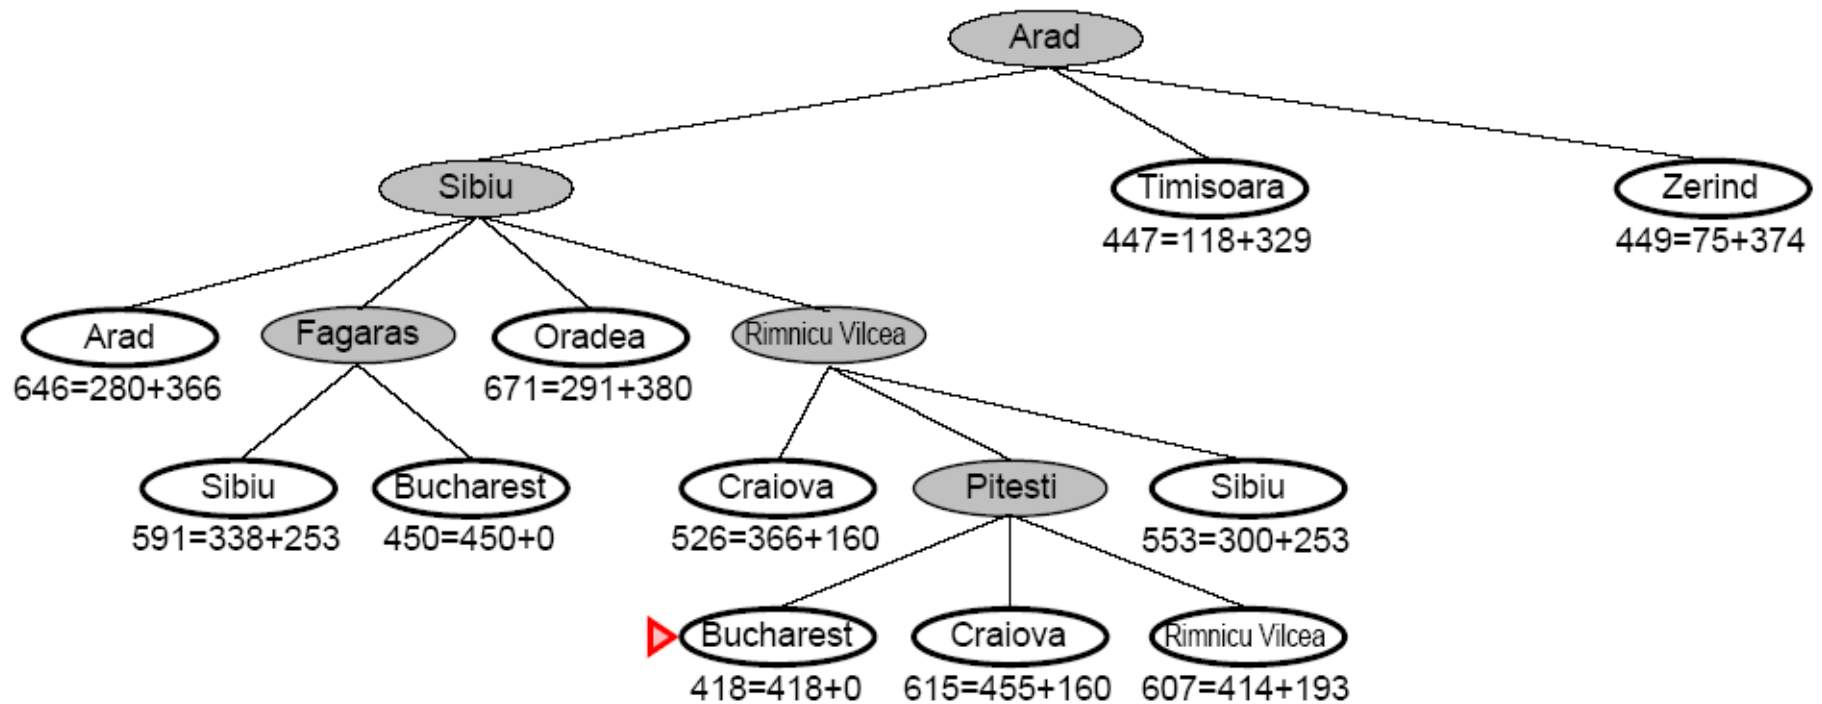
A* Search



A* Search




A* Search



A* Search

- A* expands nodes in increasing f value
 - Gradually adds f-contours of nodes (like breadth-first search adding layers)
 - Contour i has all nodes $f=f_i$ where $f_i < f_{i+1}$



A* Search

- Complete
 - Yes, unless there are infinitely many nodes with $f \leq f(G)$
- Time
 - Exponential in [relative error of h x length of soln]
 - The better the heuristic, the better the time
 - Best case h is perfect, $O(d)$
 - Worst case $h = 0$, $O(b^d)$ same as BFS
- Space
 - Keeps all nodes in memory and save in case of repetition
 - This is $O(b^d)$ or worse
 - A* usually runs out of space before it runs out of time
- Optimal
 - Yes, cannot expand f_{i+1} unless f_i is finished

Memory-Bounded Heuristic Search

- Iterative Deepening A* (IDA*)
 - Similar to Iterative Deepening Search, but cut off at $(g(n)+h(n)) > \max$ instead of $\text{depth} > \max$
 - At each iteration, cutoff is the first f-cost that exceeds the cost of the node at the previous iteration
- RBFS – see text figures 4.5 and 4.6
- Simple Memory Bounded A* (SMA*)
 - Set max to some memory bound
 - If the memory is full, to add a node drop the worst $(g+h)$ node that is already stored
 - Expands newest best leaf, deletes oldest worst leaf

Heuristic Functions

- Example: 8-Puzzle
 - Average solution cost for a random puzzle is 22 moves
 - Branching factor is about 3
 - Empty tile in the middle -> four moves
 - Empty tile on the edge -> three moves
 - Empty tile in corner -> two moves
 - 3^{22} is approx $3.1e10$
 - Get rid of repeated states
 - 181440 distinct states

7	2	4
5		6
8	3	1

Start State

1	2	3
4	5	6
7	8	

Goal State

Heuristic Functions

- To use A^* a heuristic function must be used that never overestimates the number of steps to the goal
- h_1 =the number of misplaced tiles
- h_2 =the sum of the Manhattan distances of the tiles from their goal positions

Heuristic Functions

- $h1 = 7$
- $h2 = 4+0+3+3+1+0+2+1 = 14$

7	2	4
5		6
8	3	1

Start State

1	2	3
4	5	6
7	8	

Goal State

Dominance

- If $h_2(n) > h_1(n)$ for all n (both admissible) then $h_2(n)$ dominates $h_1(n)$ and is better for the search
- Take a look at figure 4.8!

Relaxed Problems

- A Relaxed Problem is a problem with fewer restrictions on the actions
 - The cost of an optimal solution to a relaxed problem is an admissible heuristic for the original problem
- Key point: The optimal solution of a relaxed problem is no greater than the optimal solution of the real problem

Relaxed Problems

- Example: 8-puzzle
 - Consider only getting tiles 1, 2, 3, and 4 into place
 - If the rules are relaxed such that a tile can move anywhere then $h_1(n)$ gives the shortest solution
 - If the rules are relaxed such that a tile can move to any adjacent square then $h_2(n)$ gives the shortest solution

Relaxed Problems

- Store sub-problem solutions in a database
 - # patterns is much smaller than the search space
 - Generate database by working backwards from the solution
 - If multiple sub-problems apply, take the max
 - If multiple disjoint sub-problems apply, heuristics can be added

Learning Heuristics From Experience

- $h(n)$ is an estimate cost of the solution beginning at state n
- How can an agent construct such a function?
- Experience!
 - Have the agent solve many instances of the problem and store the actual cost of $h(n)$ at some state n
 - Learn from the features of a state that are relevant to the solution, rather than the state itself
 - Generate “many” states with a given feature and determine the average distance
 - Combine the information from multiple features
 - $h(n) = c(1)*x_1(n) + c(2)*x_2(n) + \dots$ where x_1, x_2, \dots are features

Optimization Problems

- Instead of considering the whole state space, consider only the current state
- Limits necessary memory; paths not retained
- Amenable to large or continuous (infinite) state spaces where exhaustive search algorithms are not possible
- Local search algorithms can't backtrack

Local Search Algorithms


- They are useful for solving **optimization problems**
 - Aim is to find a best state according to an **objective function**
- Many optimization problems do not fit the standard search model outlined in chapter 3
 - E.g. There is no goal test or path cost in Darwinian evolution
- State space landscape

Optimization Problems

- Given measure of goodness (of fit)
 - Find optimal parameters (e.g correspondences)
 - That maximize goodness measure (or minimize badness measure)
- Optimization techniques
 - Direct (closed-form)
 - Search (generate-test)
 - Heuristic search (e.g Hill Climbing)
 - Genetic Algorithm

Direct Optimization

- The slope of a function at the maximum or minimum is 0
 - Function is neither growing nor shrinking
 - True at global, but also local extreme points
- Find where the slope is zero and you find extrema!
- (If you have the equation, use calculus (first derivative=0))



Hill Climbing

- Consider all possible successors as “one step” from the current state on the landscape.
- At each iteration, go to
 - The best successor (steepest ascent)
 - Any uphill move (first choice)
 - Any uphill move but steeper is more probable (stochastic)
- All variations get stuck at local maxima

Hill Climbing


“Like climbing Everest in thick fog with amnesia”

```
function HILL-CLIMBING(problem) returns a state that is a local maximum
  inputs: problem, a problem
  local variables: current, a node
                     neighbor, a node

  current ← MAKE-NODE(INITIAL-STATE[problem])
  loop do
    neighbor ← a highest-valued successor of current
    if VALUE[neighbor] < VALUE[current] then return STATE[current]
    current ← neighbor
  end
```


Hill Climbing

Problem: depending on initial state, can get stuck on local maxima



In continuous spaces, problems w/ choosing step size, slow convergence

Hill Climbing

- Local maxima = no uphill step
 - Algorithms on previous slide fail (not complete)
 - Allow “random restart” which is complete, but might take a very long time
- Plateau = all steps equal (flat or shoulder)
 - Must move to equal state to make progress, but no indication of the correct direction
- Ridge = narrow path of maxima, but might have to go down to go up (e.g. diagonal ridge in 4-direction space)

Simulated Annealing

- Idea: Escape local maxima by allowing some “bad” moves
 - But gradually decreasing their frequency
- Algorithm is randomized:
 - Take a step if random number is less than a value based on both the objective function and the Temperature
- When Temperature is high, chance of going toward a higher value of optimization function $J(x)$ is greater
- Note higher dimension: “perturb parameter vector” vs. “look at next and previous value”

Simulated Annealing

```
function SIMULATED-ANNEALING(problem, schedule) returns a solution state
  inputs: problem, a problem
           schedule, a mapping from time to “temperature”
  local variables: current, a node
                    next, a node
                    T, a “temperature” controlling prob. of downward steps

  current  $\leftarrow$  MAKE-NODE(INITIAL-STATE[problem])
  for t  $\leftarrow$  1 to  $\infty$  do
    T  $\leftarrow$  schedule[t]
    if T = 0 then return current
    next  $\leftarrow$  a randomly selected successor of current
     $\Delta E \leftarrow$  VALUE[next] – VALUE[current]
    if  $\Delta E > 0$  then current  $\leftarrow$  next
    else current  $\leftarrow$  next only with probability  $e^{\Delta E/T}$ 
```

Genetic Algorithms

- Quicker but randomized searching for an optimal parameter vector
- Operations
 - Crossover (2 parents -> 2 children)
 - Mutation (one bit)
- Basic structure
 - Create population
 - Perform crossover & mutation (on fittest)
 - Keep only fittest children

Genetic Algorithms

- Children carry parts of their parents' data
- Only “good” parents can reproduce
 - Children are at least as “good” as parents?
 - No, but “worse” children don’t last long
- Large population allows many “current points” in search
 - Can consider several regions (watersheds) at once

Genetic Algorithms

- Representation
 - Children (after crossover) should be similar to parent, not random
 - Binary representation of numbers isn't good - what happens when you crossover in the middle of a number?
 - Need “reasonable” breakpoints for crossover (e.g. between R, xcenter and ycenter but not within them)
- “Cover”
 - Population should be large enough to “cover” the range of possibilities
 - Information shouldn't be lost too soon
 - Mutation helps with this issue

Experimenting With GAs

- Be sure you have a reasonable “goodness” criterion
- Choose a good representation (including methods for crossover and mutation)
- Generate a sufficiently random, large enough population
- Run the algorithm “long enough”
- Find the “winners” among the population
- Variations: multiple populations, keeping vs. not keeping parents, “immigration / emigration”, mutation rate, etc.

Iterative Deepening Search

```
function ITERATIVE-DEEPENING-SEARCH(problem) returns a solution
  inputs: problem, a problem
  for depth  $\leftarrow$  0 to  $\infty$  do
    result  $\leftarrow$  DEPTH-LIMITED-SEARCH(problem, depth)
    if result  $\neq$  cutoff then return result
  end
```




Iterative Deepening Search

Limit = 0




Iterative Deepening Search

Limit = 1




Iterative Deepening Search



Iterative Deepening Search

Limit = 3



Iterative Deepening Search

- Complete
 - Yes
- Time
 - $O(b^d)$
- Space
 - $O(bd)$
- Optimal
 - Yes if step cost = 1
 - Can be modified to explore uniform cost tree


Lessons From Iterative Deepening Search

- Faster than BFS even though IDS generates repeated states
 - BFS generates nodes up to level $d+1$
 - IDS only generates nodes up to level d
- In general, iterative deepening search is the preferred uninformed search method when there is a large search space and the depth of the solution is not known

Avoiding Repeated States

- Complication of wasting time by expanding states that have already been encountered and expanded before
 - Failure to detect repeated states can turn a linear problem into an exponential one
- Sometimes, repeated states are unavoidable
 - Problems where the actions are reversible
 - Route finding
 - Sliding blocks puzzles

Avoiding Repeated States



State Space

Search Tree

Avoiding Repeated States

```
function GRAPH-SEARCH( problem, fringe) returns a solution, or failure
  closed  $\leftarrow$  an empty set
  fringe  $\leftarrow$  INSERT(MAKE-NODE(INITIAL-STATE[problem]), fringe)
  loop do
    if fringe is empty then return failure
    node  $\leftarrow$  REMOVE-FRONT(fringe)
    if GOAL-TEST[problem](STATE[node]) then return node
    if STATE[node] is not in closed then
      add STATE[node] to closed
      fringe  $\leftarrow$  INSERTALL(EXPAND(node, problem), fringe)
  end
```

THANKS

金 融 先 锋 科 技 向 善